

Dataset Link: <https://microdata.statistics.gov.rw/index.php/catalog/82>

Github Link: <https://github.com/boaziza/formative2-pca-pair-team-34.git>

Formative Assignment: Advanced Linear Algebra (PCA)

This notebook will guide you through the implementation of Principal Component Analysis (PCA). Fill in the missing code and provide the required answers in the appropriate sections. You will work with a dataset that is Africanized .

1. Make sure to display outputs for each code cell when submitting.
2. Do not write all your code on one cell
3. Do not use any libraries aside from numpy

```
In [1]: # Setup: only numpy and matplotlib are used in this notebook
import numpy as np
import matplotlib.pyplot as plt
import time # standard-library timer, used only for benchmarking in Task 3

np.set_printoptions(precision=4, suppress=True, linewidth=120)

# Path to the dataset (CSV exported from NISR EICV5 microdata)
DATA_PATH = "data/eicv5_pca_subset.csv"

print("numpy", np.__version__, "| matplotlib", plt.matplotlib.__version__)
print("Data file:", DATA_PATH)
```

```
numpy 2.5.3 | matplotlib 3.11.2
Data file: data/eicv5_pca_subset.csv
```

Data Handling: load, inspect missing values, impute and encode

Non-numeric columns are encoded and missing values imputed before standardization.

```
In [2]: # Load the raw CSV with numpy only: read every cell as text first,
# then decide column by column which ones are numeric.
raw = np.genfromtxt(DATA_PATH, delimiter=",", dtype=str, encoding="utf-8")
raw = np.char.strip(np.char.strip(raw), '"') # remove spaces and quote marks

column_names = [str(c) for c in raw[0]]
data_str = raw[1:]
print("Shape (rows, columns):", data_str.shape)
print("Columns:", list(column_names))
data_str[:5]
```

```
Shape (rows, columns): (14580, 13)
Columns: ['province', 'urban_rural', 'hh_size', 'consumption_per_ae', 'food_share',
'food_expenditure', 'own_food_production', 'annual_nonfood_exp', 'education_exp', 'water_exp', 'electricity_exp', 'durable_goods_value', 'transfers_received']
```

```
Out[2]: array([[ 'Kigali City', 'Urban', '2.0000', '666627.3125', '0.5264', '665029.9989',
'19770.8340', '35500.0000',
'0.0000', '12000.0000', '13000.0000', '2555.5555', '17000.0000'],
[ 'Kigali City', 'Urban', '4.0000', '653076.2072', '0.4182', '836336.6638',
'0.0000', '40000.0000',
'33000.0000', '48000.0000', '78000.0000', '16944.4444', '16500.0000'],
[ 'Kigali City', 'Urban', '4.0000', '900828.9873', '0.3969', '1160213.3325',
'0.0000', '119700.0000', '0.0000',
'19821.4297', '13000.0000', '14000.0000', '1000.0000'],
[ 'Kigali City', 'Urban', '3.0000', '1186645.0990', '0.6319', '1421431.653
3', '0.0000', '17500.0000',
'30000.0000', '1564.2858', '0.0000', '20013.8887', '3000.0000'],
[ 'Kigali City', 'Urban', '4.0000', '1099580.5585', '0.2744', '1006791.668
0', '0.0000', '83000.0000',
'453000.0000', '75000.0000', '78000.0000', '105777.7779', '17500.0000']],
dtype='<U19')
```

```
In [3]: # Identify numeric vs non-numeric columns and count missing values
MISSING_TOKENS = {"", "nan", "na", "n/a", "null", ".", "none", "?"}

def is_missing(arr):
    return np.isin(np.char.lower(arr), list(MISSING_TOKENS))

def is_numeric_column(col):
    present = col[~is_missing(col)]
    try:
        present.astype(float)
        return True
    except ValueError:
        return False

numeric_idx = [j for j in range(data_str.shape[1]) if is_numeric_column(data_str[:, j])]
categorical_idx = [j for j in range(data_str.shape[1]) if j not in numeric_idx]

print("Numeric columns:      ", [column_names[j] for j in numeric_idx])
print("Non-numeric columns: ", [column_names[j] for j in categorical_idx])
print()
print(f"{'Column':<25}{'Type':<14}{'Missing':>8}{'% Missing':>11}")
for j, name in enumerate(column_names):
    n_miss = is_missing(data_str[:, j]).sum()
    kind = "numeric" if j in numeric_idx else "categorical"
    print(f"{'name':<25}{'kind':<14}{'n_miss':>8}{'100 * n_miss / len(data_str):>10.2f}%")
print("\nTotal missing cells:", is_missing(data_str).sum())
```

Numeric columns: ['hh_size', 'consumption_per_ae', 'food_share', 'food_expenditure', 'own_food_production', 'annual_nonfood_exp', 'education_exp', 'water_exp', 'electricity_exp', 'durable_goods_value', 'transfers_received']
 Non-numeric columns: ['province', 'urban_rural']

Column	Type	Missing	% Missing
province	categorical	0	0.00%
urban_rural	categorical	0	0.00%
hh_size	numeric	0	0.00%
consumption_per_ae	numeric	0	0.00%
food_share	numeric	0	0.00%
food_expenditure	numeric	0	0.00%
own_food_production	numeric	0	0.00%
annual_nonfood_exp	numeric	0	0.00%
education_exp	numeric	0	0.00%
water_exp	numeric	0	0.00%
electricity_exp	numeric	0	0.00%
durable_goods_value	numeric	6	0.04%
transfers_received	numeric	5	0.03%

Total missing cells: 11

```
In [4]: # Impute missing values
#   numeric columns   -> median (robust to the heavy right skew of income/consump
#   categorical columns -> most frequent category (mode)
numeric_data = np.empty((data_str.shape[0], len(numeric_idx)))
for k, j in enumerate(numeric_idx):
    col = data_str[:, j].copy()
    col[is_missing(col)] = "nan"
    numeric_data[:, k] = col.astype(float)

numeric_medians = np.nanmedian(numeric_data, axis=0)
nan_rows, nan_cols = np.where(np.isnan(numeric_data))
numeric_data[nan_rows, nan_cols] = numeric_medians[nan_cols]

categorical_data = data_str[:, categorical_idx].copy()
for k in range(categorical_data.shape[1]):
    col = categorical_data[:, k]
    miss = is_missing(col)
    values, counts = np.unique(col[~miss], return_counts=True)
    mode = values[np.argmax(counts)]
    col[miss] = mode
    print(f"{column_names[categorical_idx[k]]}: filled {miss.sum()} missing with '{

for k, j in enumerate(numeric_idx):
    print(f"{column_names[j]}: median used = {numeric_medians[k]:.3f}")

print("\nRemaining NaNs in numeric block:", np.isnan(numeric_data).sum())
```

```

province: filled 0 missing with 'Southern'
urban_rural: filled 0 missing with 'Rural'
hh_size: median used = 4.000
consumption_per_ae: median used = 256655.844
food_share: median used = 0.702
food_expenditure: median used = 399153.574
own_food_production: median used = 122501.821
annual_nonfood_exp: median used = 36300.000
education_exp: median used = 4350.000
water_exp: median used = 0.000
electricity_exp: median used = 0.000
durable_goods_value: median used = 1250.000
transfers_received: median used = 16400.000

```

Remaining NaNs in numeric block: 0

```

In [5]: # Log-transform the money columns: household spending in Rwanda is very right-skewed
# (a few households spend 100x the median), and without this the outliers would
# dominate the covariance matrix. log(1 + x) keeps zero values valid.
MONEY_COLUMNS = ["consumption_per_ae", "food_expenditure", "own_food_production", "
                 "education_exp", "water_exp", "electricity_exp", "durable_goods_value"]
numeric_names = [column_names[j] for j in numeric_idx]
for name in MONEY_COLUMNS:
    if name in numeric_names:
        k = numeric_names.index(name)
        before = numeric_data[:, k].max() / max(np.median(numeric_data[:, k]), 1)
        numeric_data[:, k] = np.log1p(numeric_data[:, k])
        print(f"{name:<22} log1p applied (max was {before:.0f}x the median)")

```

```

consumption_per_ae    log1p applied (max was 120x the median)
food_expenditure      log1p applied (max was 22x the median)
own_food_production   log1p applied (max was 39x the median)
annual_nonfood_exp    log1p applied (max was 444x the median)
education_exp         log1p applied (max was 2,690x the median)
water_exp             log1p applied (max was 1,200,000x the median)
electricity_exp       log1p applied (max was 780,000x the median)
durable_goods_value   log1p applied (max was 51,945x the median)
transfers_received    log1p applied (max was 349x the median)

```

```

In [6]: # Encode non-numeric columns
# 2 categories -> single 0/1 column (e.g. Urban/Rural)
# 3 to 10 categories -> one-hot encoding minus one reference category (e.g. Province)
# no false ordering is introduced
# more than 10 -> frequency encoding (share of households in that category)
# which avoids adding dozens of sparse columns (e.g. 30 dummies for 10 categories)
encoded_blocks, encoded_names = [], []
for k in range(categorical_data.shape[1]):
    name = column_names[categorical_idx[k]]
    col = categorical_data[:, k]
    categories, codes, counts = np.unique(col, return_inverse=True, return_counts=True)
    if len(categories) == 2:
        encoded_blocks.append(codes.reshape(-1, 1).astype(float))
        encoded_names.append(f"{name}={categories[1]}")
        print(f"{name}: binary (1 = {categories[1]}, 0 = {categories[0]})")
    elif len(categories) <= 10:
        # drop the first category (reference) so the dummies are not perfectly collinear

```

```

# which would otherwise create a zero eigenvalue
one_hot = (codes.reshape(-1, 1) == np.arange(1, len(categories))).astype(float)
encoded_blocks.append(one_hot)
encoded_names += [f"{name}={c}" for c in categories[1:]]
print(f"{name}: one-hot -> {len(categories) - 1} columns (reference = {categories[0]})")
else:
    freq = counts[codes] / len(col)
    encoded_blocks.append(freq.reshape(-1, 1))
    encoded_names.append(f"{name}_freq")
    print(f"{name}: frequency encoded ({len(categories)} categories)")

# Keep a label array for colouring the plots in Step 8 (first categorical column)
group_labels = categorical_data[:, 0] if categorical_data.shape[1] > 0 else None

X = np.hstack([numeric_data] + encoded_blocks)
feature_names = [column_names[j] for j in numeric_idx] + encoded_names
print("\nFinal numeric matrix shape:", X.shape)
print("Features:", feature_names)

```


province: one-hot -> 4 columns (reference = Eastern)
 urban_rural: binary (1 = Urban, 0 = Rural)

Final numeric matrix shape: (14580, 16)
 Features: ['hh_size', 'consumption_per_ae', 'food_share', 'food_expenditure', 'own_food_production', 'annual_nonfood_exp', 'education_exp', 'water_exp', 'electricity_exp', 'durable_goods_value', 'transfers_received', 'province=Kigali City', 'province=Northern', 'province=Southern', 'province=Western', 'urban_rural=Urban']

Step 1: Load and Standardize the Data

Before applying PCA, we must standardize the dataset. Standardization ensures that all features have a mean of 0 and a standard deviation of 1, which is essential for PCA. Fill in the code to standardize the dataset.

STRICTLY - Write code that implements standardization based on the image below



```

In [7]: # Step 1: Load and Standardize the data (use of numpy only allowed)
# Z = (X - mean) / standard deviation, computed column by column
mean = np.mean(X, axis=0)
std = np.std(X, axis=0)
std[std == 0] = 1.0 # guard against a constant column (division by zero)

standardized_data = (X - mean) / std # Do not use sklearn

```

```
print("Means after standardizing (should be ~0):", standardized_data.mean(axis=0).r
print("Std after standardizing (should be 1):  ", standardized_data.std(axis=0).ro
standardized_data[:5] # Display the first few rows of standardized data
```

```
Means after standardizing (should be ~0): [-0.  0. -0. -0. -0.  0. -0. -0.  0.  0.
0.  0.  0. -0. -0. -0.]
Std after standardizing (should be 1):  [1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.  1.
1.  1.]
```

```
Out[7]: array([[ -1.1383,  1.1382, -0.9442,  0.5075,  0.0434,  0.068 , -1.3204,  1.5017,
 1.7056,  0.4031,  0.221 ,  2.8284,
 -0.4439, -0.5979, -0.5472,  2.1845],
 [-0.1941,  1.1113, -1.6809,  0.7207, -2.1047,  0.1462,  0.8563,  1.8202,
 2.1315,  0.8815,  0.2096,  2.8284,
 -0.4439, -0.5979, -0.5472,  2.1845],
 [-0.1941,  1.534 , -1.8259,  1.0251, -2.1047,  0.8642, -1.3204,  1.617 ,
 1.7056,  0.8332, -0.8629,  2.8284,
 -0.4439, -0.5979, -0.5472,  2.1845],
 [-0.6662,  1.8961, -0.2259,  1.214 , -2.1047, -0.3953,  0.8364,  1.0336, -
0.5459,  0.9236, -0.4427,  2.8284,
 -0.4439, -0.5979, -0.5472,  2.1845],
 [-0.1941,  1.796 , -2.6599,  0.8932, -2.1047,  0.6244,  1.4043,  1.9227,
 2.1315,  1.3448,  0.2321,  2.8284,
 -0.4439, -0.5979, -0.5472,  2.1845]])
```

Step 3: Calculate the Covariance Matrix

The covariance matrix helps us understand how the features are related to each other. It is a key component in PCA.

```
In [8]: # Step 3: Calculate the Covariance Matrix
# Cov = (Z^T Z) / (n - 1); Z is already centred, so no extra mean subtraction is ne
n_samples = standardized_data.shape[0]
cov_matrix = (standardized_data.T @ standardized_data) / (n_samples - 1) # Calcula

# Sanity check against numpy's built-in version
print("Matches np.cov:", np.allclose(cov_matrix, np.cov(standardized_data, rowvar=F
print("Shape:", cov_matrix.shape)
cov_matrix
```

```
Matches np.cov: True
Shape: (16, 16)
```

```

Out[8]: array([[ 1.0001, -0.289 , -0.0031,  0.3035,  0.2531,  0.2981,  0.615 ,  0.0187,
 0.0787,  0.2278,  0.0177, -0.0654,
-0.0093, -0.0075,  0.066 , -0.0355],
 [ -0.289 ,  1.0001, -0.5696,  0.4592, -0.327 ,  0.4089, -0.056 ,  0.4121,
 0.5199,  0.485 , -0.0156,  0.4449,
-0.0922, -0.0969, -0.1376,  0.4395],
 [ -0.0031, -0.5696,  1.0001, -0.2244,  0.4383, -0.3999, -0.194 , -0.4598, -
 0.5943, -0.5199,  0.1126, -0.4688,
 0.0987,  0.1109,  0.1239, -0.5403],
 [ 0.3035,  0.4592, -0.2244,  1.0001, -0.1318,  0.4587,  0.2604,  0.2785,
 0.3495,  0.3948, -0.0387,  0.2384,
-0.1419, -0.0238, -0.0451,  0.2559],
 [ 0.2531, -0.327 ,  0.4383, -0.1318,  1.0001, -0.005 ,  0.1756, -0.3168, -
 0.3671, -0.0909,  0.1465, -0.5054,
 0.1217,  0.0897,  0.0729, -0.5202],
 [ 0.2981,  0.4089, -0.3999,  0.4587, -0.005 ,  1.0001,  0.2976,  0.2534,
 0.3203,  0.4818,  0.0061,  0.1644,
 0.0032, -0.087 , -0.0565,  0.178 ],
 [ 0.615 , -0.056 , -0.194 ,  0.2604,  0.1756,  0.2976,  1.0001,  0.0889,
 0.1574,  0.2626,  0.0063,  0.0139,
-0.0052, -0.0257,  0.0165,  0.0507],
 [ 0.0187,  0.4121, -0.4598,  0.2785, -0.3168,  0.2534,  0.0889,  1.0001,
 0.4391,  0.3149, -0.0954,  0.348 ,
-0.0448, -0.1647, -0.127 ,  0.3949],
 [ 0.0787,  0.5199, -0.5943,  0.3495, -0.3671,  0.3203,  0.1574,  0.4391,
 1.0001,  0.4634, -0.1077,  0.4672,
-0.0744, -0.1408, -0.0559,  0.517 ],
 [ 0.2278,  0.485 , -0.5199,  0.3948, -0.0909,  0.4818,  0.2626,  0.3149,
 0.4634,  1.0001, -0.0517,  0.2556,
-0.0316, -0.1128, -0.0864,  0.2982],
 [ 0.0177, -0.0156,  0.1126, -0.0387,  0.1465,  0.0061,  0.0063, -0.0954, -
 0.1077, -0.0517,  1.0001, -0.1056,
 0.056 , -0.0092,  0.0705, -0.1834],
 [ -0.0654,  0.4449, -0.4688,  0.2384, -0.5054,  0.1644,  0.0139,  0.348 ,
 0.4672,  0.2556, -0.1056,  1.0001,
-0.157 , -0.2114, -0.1935,  0.5544],
 [ -0.0093, -0.0922,  0.0987, -0.1419,  0.1217,  0.0032, -0.0052, -0.0448, -
 0.0744, -0.0316,  0.056 , -0.157 ,
 1.0001, -0.2654, -0.2429, -0.0859],
 [ -0.0075, -0.0969,  0.1109, -0.0238,  0.0897, -0.087 , -0.0257, -0.1647, -
 0.1408, -0.1128, -0.0092, -0.2114,
-0.2654,  1.0001, -0.3272, -0.1256],
 [ 0.066 , -0.1376,  0.1239, -0.0451,  0.0729, -0.0565,  0.0165, -0.127 , -
 0.0559, -0.0864,  0.0705, -0.1935,
-0.2429, -0.3272,  1.0001, -0.0698],
 [ -0.0355,  0.4395, -0.5403,  0.2559, -0.5202,  0.178 ,  0.0507,  0.3949,
 0.517 ,  0.2982, -0.1834,  0.5544,
-0.0859, -0.1256, -0.0698,  1.0001]])

```

In a paragraph that has a maximum of 5 Lines, Explain why we need to compute a covariance matrix, provide atleast 2 reasons

Answer: We compute the covariance matrix because (1) it summarises how every pair of features moves together; for example, `consumption_per_ae` and `food_share` have a

covariance of -0.57, so richer households spend a smaller share on food, and PCA uses this to group redundant features. (2) Its eigenvectors are the directions of maximum spread in the data and its eigenvalues are the variance along each direction, so PCA is simply the eigendecomposition of this matrix. (3) Because the data is standardized, the covariance matrix equals the correlation matrix, so features measured in RWF and features that are ratios or 0/1 flags all get equal weight.

Step 4: Perform Eigendecomposition

Eigendecomposition of the covariance matrix will give us the eigenvalues and eigenvectors, which are essential for PCA. Fill in the code to compute the eigenvalues and eigenvectors of the covariance matrix.

```
In [9]: # Step 4: Perform Eigendecomposition
# The covariance matrix is symmetric, so eig is used: it guarantees real eigenvalues
# and orthonormal eigenvectors (np.linalg.eig can return tiny complex parts)
eigenvalues, eigenvectors = np.linalg.eigh(cov_matrix) # Perform eigendecomposition

# Check: C v = lambda v for every pair
print("C v = lambda v holds:", np.allclose(cov_matrix @ eigenvectors, eigenvectors
eigenvalues, eigenvectors
```

C v = lambda v holds: True


```

Out[9]: (array([0.1731, 0.2521, 0.2714, 0.4211, 0.4342, 0.4583, 0.5257, 0.5521, 0.6636, 0.
7429, 0.8873, 1.1352, 1.2746, 1.3482,
2.2365, 4.6249])),
array([[ -0.5295, -0.2136, 0.2386, -0.028 , -0.1589, 0.3526, -0.0033, 0.0394,
0.009 , -0.1202, 0.1792, 0.3098,
0.1298, -0.0217, -0.5516, -0.0319],
[ -0.6791, -0.1183, -0.0312, -0.0141, 0.1372, -0.3661, -0.1191, -0.0335, -
0.0538, -0.0288, -0.1669, -0.4105,
-0.1182, 0.0731, 0.1197, -0.3489],
[ -0.2183, 0.332 , -0.618 , -0.0991, -0.1343, 0.1 , -0.2257, -0.033 ,
0.0626, -0.4698, -0.0633, -0.0339,
0.0232, -0.01 , -0.0316, 0.3726],
[ 0.3552, -0.0789, 0.2387, 0.067 , 0.1896, -0.0398, -0.1837, 0.2137,
0.0241, -0.6841, -0.1441, -0.213 ,
0.1009, 0.1117, -0.2569, -0.2566],
[ 0.1676, -0.1418, 0.1845, -0.359 , -0.3657, -0.2965, -0.361 , -0.3581,
0.0763, 0.1569, -0.1921, -0.1724,
-0.1423, 0.041 , -0.3572, 0.249 ],
[ 0.0822, 0.159 , -0.199 , -0.0227, -0.3833, 0.0207, 0.6886, -0.0263, -
0.1211, -0.0294, -0.2021, -0.2459,
-0.1345, 0.0441, -0.326 , -0.2501],
[ 0.0664, 0.1759, -0.2938, 0.1133, 0.3669, -0.5573, 0.0084, -0.0367, -
0.0203, 0.1365, 0.2265, 0.2941,
0.0764, -0.0071, -0.4966, -0.0977],
[ 0.0306, -0.076 , -0.137 , -0.1067, -0.0012, 0.0573, 0.0586, -0.093 ,
0.9204, 0.0013, 0.0268, 0.0683,
-0.0753, -0.0411, 0.0388, -0.2939],
[ 0.0588, 0.2011, -0.0774, 0.4697, -0.5534, -0.0274, -0.446 , 0.239 , -
0.0088, 0.1742, 0.0516, 0.0409,
0.0202, -0.053 , -0.0008, -0.3584],
[ 0.1133, 0.0748, -0.2582, -0.1631, 0.3619, 0.5527, -0.2984, -0.1247, -
0.1613, 0.3233, -0.181 , -0.1576,
-0.1146, 0.0304, -0.2302, -0.3036],
[ 0.0454, 0.0567, 0.0264, -0.0609, -0.0017, 0.0711, -0.0109, 0.0731,
0.0313, 0.0105, 0.7913, -0.5482,
-0.1624, -0.108 , -0.0906, 0.0759],
[ 0.132 , -0.422 , -0.265 , 0.0125, -0.1506, 0.0004, -0.0256, -0.5765, -
0.2593, -0.2595, 0.2738, 0.1667,
0.0609, -0.0379, 0.1929, -0.3143],
[ 0.0495, -0.3729, -0.1753, -0.009 , 0.0103, -0.0411, -0.0322, 0.368 , -
0.0823, -0.0807, -0.0367, 0.2189,
-0.7698, -0.1716, -0.0319, 0.0622],
[ 0.054 , -0.389 , -0.2905, -0.0695, -0.0877, -0.0063, 0.0022, 0.3162,
0.028 , 0.1594, 0.0915, -0.0599,
0.2088, 0.7476, -0.0031, 0.0906],
[ 0.0557, -0.4143, -0.2536, -0.0589, -0.0196, -0.05 , 0.0216, 0.2681,
0.0059, 0.1275, -0.1452, -0.2183,
0.4674, -0.6065, -0.0748, 0.0724],
[ 0.0166, 0.229 , 0.0372, -0.756 , -0.1367, -0.1114, -0.0544, 0.3065, -
0.1505, -0.0377, 0.1347, 0.2317,
0.1019, -0.0563, 0.1652, -0.3311]]))

```

Step 5: Sort Principal Components

Sort the eigenvectors based on their corresponding eigenvalues in descending order. The higher the eigenvalue, the more important the eigenvector. Complete the code to sort the eigenvectors and print the sorted components.

How Is Explained Variance Used In PCA?

```
In [10]: # Step 5: Sort Principal Components
sorted_indices = np.argsort(eigenvalues)[::-1] # Sort eigenvalues in descending order
sorted_eigenvalues = eigenvalues[sorted_indices]
sorted_eigenvectors = eigenvectors[:, sorted_indices] # Sort eigenvectors accordingly

print("Sorted eigenvalues:", sorted_eigenvalues)
sorted_eigenvectors
```

```
Sorted eigenvalues: [4.6249 2.2365 1.3482 1.2746 1.1352 0.8873 0.7429 0.6636 0.5521
0.5257 0.4583 0.4342 0.4211 0.2714 0.2521 0.1731]
```

```

Out[10]: array([[ -0.0319, -0.5516, -0.0217,  0.1298,  0.3098,  0.1792, -0.1202,  0.009 ,
  0.0394, -0.0033,  0.3526, -0.1589,
        -0.028 ,  0.2386, -0.2136, -0.5295],
 [ -0.3489,  0.1197,  0.0731, -0.1182, -0.4105, -0.1669, -0.0288, -0.0538, -
 0.0335, -0.1191, -0.3661,  0.1372,
        -0.0141, -0.0312, -0.1183, -0.6791],
 [  0.3726, -0.0316, -0.01  ,  0.0232, -0.0339, -0.0633, -0.4698,  0.0626, -
 0.033 , -0.2257,  0.1   , -0.1343,
        -0.0991, -0.618 ,  0.332 , -0.2183],
 [ -0.2566, -0.2569,  0.1117,  0.1009, -0.213 , -0.1441, -0.6841,  0.0241,
 0.2137, -0.1837, -0.0398,  0.1896,
        0.067 ,  0.2387, -0.0789,  0.3552],
 [  0.249 , -0.3572,  0.041 , -0.1423, -0.1724, -0.1921,  0.1569,  0.0763, -
 0.3581, -0.361 , -0.2965, -0.3657,
        -0.359 ,  0.1845, -0.1418,  0.1676],
 [ -0.2501, -0.326 ,  0.0441, -0.1345, -0.2459, -0.2021, -0.0294, -0.1211, -
 0.0263,  0.6886,  0.0207, -0.3833,
        -0.0227, -0.199 ,  0.159 ,  0.0822],
 [ -0.0977, -0.4966, -0.0071,  0.0764,  0.2941,  0.2265,  0.1365, -0.0203, -
 0.0367,  0.0084, -0.5573,  0.3669,
        0.1133, -0.2938,  0.1759,  0.0664],
 [ -0.2939,  0.0388, -0.0411, -0.0753,  0.0683,  0.0268,  0.0013,  0.9204, -
 0.093 ,  0.0586,  0.0573, -0.0012,
        -0.1067, -0.137 , -0.076 ,  0.0306],
 [ -0.3584, -0.0008, -0.053 ,  0.0202,  0.0409,  0.0516,  0.1742, -0.0088,
 0.239 , -0.446 , -0.0274, -0.5534,
        0.4697, -0.0774,  0.2011,  0.0588],
 [ -0.3036, -0.2302,  0.0304, -0.1146, -0.1576, -0.181 ,  0.3233, -0.1613, -
 0.1247, -0.2984,  0.5527,  0.3619,
        -0.1631, -0.2582,  0.0748,  0.1133],
 [  0.0759, -0.0906, -0.108 , -0.1624, -0.5482,  0.7913,  0.0105,  0.0313,
 0.0731, -0.0109,  0.0711, -0.0017,
        -0.0609,  0.0264,  0.0567,  0.0454],
 [ -0.3143,  0.1929, -0.0379,  0.0609,  0.1667,  0.2738, -0.2595, -0.2593, -
 0.5765, -0.0256,  0.0004, -0.1506,
        0.0125, -0.265 , -0.422 ,  0.132 ],
 [  0.0622, -0.0319, -0.1716, -0.7698,  0.2189, -0.0367, -0.0807, -0.0823,
 0.368 , -0.0322, -0.0411,  0.0103,
        -0.009 , -0.1753, -0.3729,  0.0495],
 [  0.0906, -0.0031,  0.7476,  0.2088, -0.0599,  0.0915,  0.1594,  0.028 ,
 0.3162,  0.0022, -0.0063, -0.0877,
        -0.0695, -0.2905, -0.389 ,  0.054 ],
 [  0.0724, -0.0748, -0.6065,  0.4674, -0.2183, -0.1452,  0.1275,  0.0059,
 0.2681,  0.0216, -0.05  , -0.0196,
        -0.0589, -0.2536, -0.4143,  0.0557],
 [ -0.3311,  0.1652, -0.0563,  0.1019,  0.2317,  0.1347, -0.0377, -0.1505,
 0.3065, -0.0544, -0.1114, -0.1367,
        -0.756 ,  0.0372,  0.229 ,  0.0166]])

```

Task 2: Explained Variance and Dynamic Selection of Components

Each eigenvalue is the variance captured by its principal component. Dividing by the sum of all eigenvalues gives the share of total variance explained. The number of components is

chosen automatically as the smallest k whose cumulative explained variance reaches the threshold below.

```
In [11]: # Task 2: Explained variance ratio and cumulative explained variance
explained_variance_ratio = sorted_eigenvalues / np.sum(sorted_eigenvalues)
cumulative_variance = np.cumsum(explained_variance_ratio)

print(f"{'PC':<6}{'Eigenvalue':>12}{'Explained %':>14}{'Cumulative %':>15}")
for i, (ev, r, c) in enumerate(zip(sorted_eigenvalues, explained_variance_ratio, cu
    print(f"PC{i:<4}{ev:>12.4f}{100 * r:>13.2f}%{100 * c:>14.2f}%")
```

PC	Eigenvalue	Explained %	Cumulative %
PC1	4.6249	28.90%	28.90%
PC2	2.2365	13.98%	42.88%
PC3	1.3482	8.43%	51.31%
PC4	1.2746	7.97%	59.27%
PC5	1.1352	7.09%	66.37%
PC6	0.8873	5.55%	71.91%
PC7	0.7429	4.64%	76.55%
PC8	0.6636	4.15%	80.70%
PC9	0.5521	3.45%	84.15%
PC10	0.5257	3.29%	87.44%
PC11	0.4583	2.86%	90.30%
PC12	0.4342	2.71%	93.02%
PC13	0.4211	2.63%	95.65%
PC14	0.2714	1.70%	97.34%
PC15	0.2521	1.58%	98.92%
PC16	0.1731	1.08%	100.00%

```
In [12]: # Dynamic selection: smallest k whose cumulative explained variance >= threshold
VARIANCE_THRESHOLD = 0.80 # set by the team; justify this choice in the answer be

def select_num_components(eigvals_sorted, threshold):
    cumulative = np.cumsum(eigvals_sorted) / np.sum(eigvals_sorted)
    return int(np.argmax(cumulative >= threshold) + 1)

k_selected = select_num_components(sorted_eigenvalues, VARIANCE_THRESHOLD)
k_kaiser = int(np.sum(sorted_eigenvalues > 1)) # Kaiser rule, shown for compariso

print(f"Components needed for {VARIANCE_THRESHOLD:.0%} of variance: {k_selected} "
      f"(retains {cumulative_variance[k_selected - 1]:.2%}, discards {1 - cumulativ
print(f"Kaiser rule (eigenvalue > 1) would keep: {k_kaiser}")
print(f"Dimensions reduced from {standardized_data.shape[1]} to {k_selected}")


pcs = np.arange(1, len(sorted_eigenvalues) + 1)
fig, ax = plt.subplots(1, 2, figsize=(13, 4.5))
ax[0].bar(pcs, 100 * explained_variance_ratio, color="steelblue")
ax[0].axhline(100 / len(pcs), color="grey", ls=":", label="Average share per featur
ax[0].set(title="Scree plot", xlabel="Principal component", ylabel="Explained varia
ax[0].legend()
ax[1].plot(pcs, 100 * cumulative_variance, "o-", color="darkorange")
ax[1].axhline(100 * VARIANCE_THRESHOLD, color="red", ls="--", label=f"{VARIANCE_THR
ax[1].axvline(k_selected, color="green", ls="--", label=f"k = {k_selected}")
ax[1].set(title="Cumulative explained variance", xlabel="Number of components", yla
ax[1].legend()
```

```
plt.tight_layout()
plt.show()
```

Components needed for 80% of variance: 8 (retains 80.70%, discards 19.30%)

Kaiser rule (eigenvalue > 1) would keep: 5

Dimensions reduced from 16 to 8



```
In [13]: # Loadings: which original features drive each retained component
# (use this table when explaining what each PC represents and what is lost)
print(f'{"Feature":<28}" + "'.join(f'{"PC" + str(i + 1):>9}" for i in range(k_selected))
for f, row in zip(feature_names, sorted_eigenvectors[:, :k_selected]):
    print(f'{"f":<28}" + "'.join(f'{"v:>9.3f}" for v in row))

# How much of each original feature's variance survives in the kept components
retained_per_feature = (sorted_eigenvectors[:, :k_selected] ** 2 * sorted_eigenvalues).sum(axis=1)
print("\nShare of each feature's variance kept after reduction:")
for f, r in sorted(zip(feature_names, retained_per_feature), key=lambda t: t[1]):
    print(f'{"f":<28}{"r:>7.1%"}')
```

Feature	PC1	PC2	PC3	PC4	PC5	PC6
PC7 PC8						
hh_size	-0.032	-0.552	-0.022	0.130	0.310	0.179
-0.120 0.009						
consumption_per_ae	-0.349	0.120	0.073	-0.118	-0.411	-0.167
-0.029 -0.054						
food_share	0.373	-0.032	-0.010	0.023	-0.034	-0.063
-0.470 0.063						
food_expenditure	-0.257	-0.257	0.112	0.101	-0.213	-0.144
-0.684 0.024						
own_food_production	0.249	-0.357	0.041	-0.142	-0.172	-0.192
0.157 0.076						
annual_nonfood_exp	-0.250	-0.326	0.044	-0.135	-0.246	-0.202
-0.029 -0.121						
education_exp	-0.098	-0.497	-0.007	0.076	0.294	0.226
0.136 -0.020						
water_exp	-0.294	0.039	-0.041	-0.075	0.068	0.027
0.001 0.920						
electricity_exp	-0.358	-0.001	-0.053	0.020	0.041	0.052
0.174 -0.009						
durable_goods_value	-0.304	-0.230	0.030	-0.115	-0.158	-0.181
0.323 -0.161						
transfers_received	0.076	-0.091	-0.108	-0.162	-0.548	0.791
0.010 0.031						
province=Kigali City	-0.314	0.193	-0.038	0.061	0.167	0.274
-0.259 -0.259						
province=Northern	0.062	-0.032	-0.172	-0.770	0.219	-0.037
-0.081 -0.082						
province=Southern	0.091	-0.003	0.748	0.209	-0.060	0.091
0.159 0.028						
province=Western	0.072	-0.075	-0.607	0.467	-0.218	-0.145
0.128 0.006						
urban_rural=Urban	-0.331	0.165	-0.056	0.102	0.232	0.135
-0.038 -0.150						

Share of each feature's variance kept after reduction:

electricity_exp	62.5%
annual_nonfood_exp	66.8%
urban_rural=Urban	67.9%
own_food_production	68.9%
durable_goods_value	71.5%
province=Kigali City	73.9%
education_exp	76.1%
food_share	81.6%
consumption_per_ae	83.9%
hh_size	85.5%
province=Southern	87.8%
province=Northern	88.0%
province=Western	89.6%
food_expenditure	90.0%
water_exp	98.0%
transfers_received	99.2%

Step 6: Project Data onto Principal Components

Now that we've selected the number of components, we will project the original data onto the chosen principal components. Fill in the code to perform the projection.

```
In [14]: # Step 6: Project Data onto Principal Components
num_components = k_selected # Decide on the number of principal components to keep
W = sorted_eigenvectors[:, :num_components] # projection matrix (features)
reduced_data = standardized_data @ W # Project data onto the principal components
reduced_data[:5]
```

```
Out[14]: array([[ -3.6118,  2.1892, -0.2446, -0.1741, -0.2852,  0.6113, -0.4887,  0.1672],
                [-5.1212,  1.1975, -0.3559,  0.3491,  0.9595,  1.6257, -0.2143,  0.1305],
                [-5.2221,  2.1222, -0.0972,  0.2499,  0.4734,  0.0151, -0.786 , -0.1458],
                [-3.6986,  1.5867, -0.0255,  0.4561,  0.6515,  0.6689, -1.6486, -0.4746],
                [-6.1114,  0.7333, -0.2522,  0.1758,  0.6401,  1.5126,  0.3188, -0.0121]])
```

Step 7: Output the Reduced Data

Finally, display the reduced data obtained by projecting the original dataset onto the selected principal components.

```
In [15]: # Step 7: Output the Reduced Data
print(f'Reduced Data Shape: {reduced_data.shape}') # Display reduced data shape
reduced_data[:5] # Display the first few rows of reduced data
```

Reduced Data Shape: (14580, 8)

```
Out[15]: array([[ -3.6118,  2.1892, -0.2446, -0.1741, -0.2852,  0.6113, -0.4887,  0.1672],
                [-5.1212,  1.1975, -0.3559,  0.3491,  0.9595,  1.6257, -0.2143,  0.1305],
                [-5.2221,  2.1222, -0.0972,  0.2499,  0.4734,  0.0151, -0.786 , -0.1458],
                [-3.6986,  1.5867, -0.0255,  0.4561,  0.6515,  0.6689, -1.6486, -0.4746],
                [-6.1114,  0.7333, -0.2522,  0.1758,  0.6401,  1.5126,  0.3188, -0.0121]])
```

Step 8: Visualize Before and After PCA

Now, let's plot the original data and the data after PCA to compare the reduction in dimensions visually.


```
In [16]: # Step 8: Visualize Before and After PCA
fig, axes = plt.subplots(1, 2, figsize=(14, 5.5))

def scatter_by_group(ax, x, y):
    if group_labels is None:
        ax.scatter(x, y, s=8, alpha=0.5)
        return
    for g in np.unique(group_labels):
        m = group_labels == g
        ax.scatter(x[m], y[m], s=8, alpha=0.5, label=g)
    ax.legend(title=column_names[categorical_idx[0]], fontsize=8, markerscale=2)

# Plot original data (first two features for simplicity)
scatter_by_group(axes[0], standardized_data[:, 0], standardized_data[:, 1])
axes[0].set(title="Before PCA: original feature space (standardized)",
            xlabel=feature_names[0], ylabel=feature_names[1])
```

```
# Plot reduced data after PCA
scatter_by_group(axes[1], reduced_data[:, 0], reduced_data[:, 1])
axes[1].axhline(0, color="grey", lw=0.6); axes[1].axvline(0, color="grey", lw=0.6)
axes[1].set(title="After PCA: principal component space",
            xlabel=f"PC1 ({{explained_variance_ratio[0]:.1%}} of variance)",
            ylabel=f"PC2 ({{explained_variance_ratio[1]:.1%}} of variance)")
plt.tight_layout()
plt.show()

print("Points before / after:", standardized_data.shape[0], "/", reduced_data.shape[0])
print("Variance of PC1, PC2:", reduced_data[:, :2].var(axis=0, ddof=1).round(4),
      "(equal to the top two eigenvalues:", sorted_eigenvalues[:2].round(4), ")")
print("Mean of PC scores (centred at 0):", reduced_data.mean(axis=0).round(6))
```



Points before / after: 14580 / 14580

Variance of PC1, PC2: [4.6249 2.2365] (equal to the top two eigenvalues: [4.6249 2.2365])

Mean of PC scores (centred at 0): [-0. -0. 0. -0. -0. -0. 0. -0.]

Please make sure you do not use more than 5 lines to answer each of the following points:

1. Interpret the Visual you just created of the before and after PCA
2. Explain why you selected the number of principle components, more especially explain what tradeoffs you are making
3. Clearly mention based on your use case/ dataset ("economic activity", "population pressure"), What information is lost when reducing dimensions?

1. Interpretation of the before/after plots In the graph before PCA, the normalized hh_size and consumption_per_ae variables depict overlapping distributions for each province, but with vertical bars due to the discrete nature of household size. In the plot after PCA, PC1 and PC2 are used to depict the observations, and Kigali City generally appears in negative PC1 values. PC1 has the largest spread (28.9% of the variance) and PC2 the second largest (14.0%), both are centred at 0, and all 14,580 households are still there.

2. Why we chose k components and the tradeoffs

The number of the selected principal components is 8 since this is the minimum amount of them that can be used for capturing at least 80% of variance. In total, 80.70% of the variance

is explained by the components while the rest, 19.30%, is ignored. Only PC1 and PC2 are plotted in order to visualize them, although two components explain only 42.88% of variance.

3. What information is lost when reducing dimensions (for our dataset)

The reduction in data from 16 variables to eight principal components results in a loss of about 19.30% of the total variation. This is because of small variations and other features that could assist in identifying economic performance and population stress between different houses or provinces. Moreover, the principal components are a combination of variables making them difficult to interpret.

Task 3: Optimized PCA for Large Datasets

The optimized version: (1) builds the covariance with one matrix product (BLAS) instead of Python loops, (2) uses `eigh` for the symmetric matrix, (3) can use `float32` to halve memory, and (4) can stream the data in chunks, so the full dataset never has to fit in memory at once; only running sums and a small $d \times d$ matrix are kept.

```
In [17]: # Task 3: baseline (loop-based) vs optimized PCA

def pca_naive(X, k):
    """Straightforward version: standardize, then covariance with explicit Python loops"""
    Z = (X - X.mean(axis=0)) / X.std(axis=0)
    n, d = Z.shape
    C = np.zeros((d, d))
    for i in range(d):
        for j in range(d):
            C[i, j] = np.sum(Z[:, i] * Z[:, j]) / (n - 1)
    vals, vecs = np.linalg.eig(C)
    order = np.argsort(vals.real)[::-1]
    vecs = vecs.real[:, order[:k]]
    return Z @ vecs, vals.real[order]

def pca_optimized(X, k, dtype=np.float64):
    """Vectorized version: one matrix product for the covariance, eig for the symmetric matrix"""
    X = np.asarray(X, dtype=dtype)
    mu, sd = X.mean(axis=0), X.std(axis=0)
    sd[sd == 0] = 1
    Z = (X - mu) / sd
    C = (Z.T @ Z) / (Z.shape[0] - 1)
    vals, vecs = np.linalg.eigh(C)
    order = np.argsort(vals)[::-1]
    return Z @ vecs[:, order[:k]], vals[order]

def pca_chunked(X, k, chunk_size=100_000):
    """Out-of-core version: one pass to accumulate sums, so memory use is O(d^2), n
    X can be anything that supports row slicing, e.g. np.load(..., mmap_mode='r')."""
    n, d = X.shape
    s, ss = np.zeros(d), np.zeros((d, d))
    for start in range(0, n, chunk_size):
```

```

        block = np.asarray(X[start:start + chunk_size], dtype=np.float64)
        s += block.sum(axis=0)
        ss += block.T @ block
    mu = s / n
    cov_raw = (ss - n * np.outer(mu, mu)) / (n - 1)
    sd = np.sqrt(np.diag(cov_raw) * (n - 1) / n)
    sd[sd == 0] = 1
    C = cov_raw / np.outer(sd, sd)  # covariance of the standardized data
    vals, vecs = np.linalg.eigh(C)
    order = np.argsort(vals)[::-1]
    W = vecs[:, order[k:]]
    out = np.empty((n, k))
    for start in range(0, n, chunk_size):
        block = np.asarray(X[start:start + chunk_size], dtype=np.float64)
        out[start:start + chunk_size] = ((block - mu) / sd) @ W
    return out, vals[order]

print("Defined: pca_naive, pca_optimized, pca_chunked")

```

Defined: pca_naive, pca_optimized, pca_chunked

```

In [18]: # Correctness: all three versions must agree with Steps 1-6
# (eigenvector signs are arbitrary, so compare absolute values)
r_naive, v_naive = pca_naive(X, num_components)
r_opt, v_opt = pca_optimized(X, num_components)
r_chunk, v_chunk = pca_chunked(X, num_components, chunk_size=500)

print("Naive eigenvalues match: ", np.allclose(v_naive, sorted_eigenvalues))
print("Optimized eigenvalues match:", np.allclose(v_opt, sorted_eigenvalues))
print("Chunked eigenvalues match: ", np.allclose(v_chunk, sorted_eigenvalues))
print("Projections match (naive, optimized, chunked):",
      np.allclose(np.abs(r_naive), np.abs(reduced_data)),
      np.allclose(np.abs(r_opt), np.abs(reduced_data)),
      np.allclose(np.abs(r_chunk), np.abs(reduced_data)))

```

Naive eigenvalues match: True
 Optimized eigenvalues match: True
 Chunked eigenvalues match: True
 Projections match (naive, optimized, chunked): True True True

```

In [19]: # Benchmark on larger versions of our dataset (rows resampled with small noise)
def time_it(fn, *args, repeats=3, **kw):
    best = np.inf
    for _ in range(repeats):
        t0 = time.perf_counter()
        fn(*args, **kw)
        best = min(best, time.perf_counter() - t0)
    return best

rng = np.random.default_rng(42)
sizes = sorted([len(X), 10_000, 50_000, 200_000, 1_000_000])
results = []
for n in sizes:
    X_big = X[rng.integers(0, len(X), n)] + rng.normal(0, 1e-3, (n, X.shape[1]))
    t_naive = time_it(pca_naive, X_big, num_components, repeats=1) if n <= 200_000
    t_opt64 = time_it(pca_optimized, X_big, num_components)
    t_opt32 = time_it(pca_optimized, X_big, num_components, dtype=np.float32)

```

```


t_chunk = time_it(pca_chunked, X_big, num_components, repeats=1)
results.append((n, t_naive, t_opt64, t_opt32, t_chunk))
del X_big

print(f"{'Rows':>10}{'Naive (s)':>12}{'Optimized f64':>15}{'Optimized f32':>15}{'Ch
for n, a, b, c, d in results:
    speed = f"{a / b:>9.1f}x" if not np.isnan(a) else f"{'-':>10}"
    a_txt = f"{a:>12.4f}" if not np.isnan(a) else f"{'skipped':>12}"
    print(f"{n:>10},{a_txt}{b:>15.4f}{c:>15.4f}{d:>10.4f}{speed}")

res = np.array(results, dtype=float)
plt.figure(figsize=(8, 4.5))
plt.plot(res[:, 0], res[:, 1], "o-", label="Naive (Python loops, eig)")
plt.plot(res[:, 0], res[:, 2], "s-", label="Optimized float64 (matmul, eigh)")
plt.plot(res[:, 0], res[:, 3], "^-", label="Optimized float32")
plt.plot(res[:, 0], res[:, 4], "d-", label="Chunked / out-of-core")
plt.xscale("log"); plt.yscale("log")
plt.xlabel("Number of rows"); plt.ylabel("Time (seconds, log scale)")
plt.title("PCA runtime vs dataset size")
plt.legend(); plt.grid(True, which="both", alpha=0.3)
plt.show()

```

Rows	Naive (s)	Optimized f64	Optimized f32	Chunked	Speed-up
10,000	0.0072	0.0018	0.0011	0.0012	3.9x
14,580	0.0125	0.0022	0.0015	0.0017	5.6x
50,000	0.0422	0.0081	0.0065	0.0067	5.2x
200,000	0.2976	0.0378	0.0276	0.0285	7.9x
1,000,000	skipped	0.2041	0.1341	0.1300	-



<i>Please ensure this tracking sheet is completed as a group, as it will be used to monitor your individual participation within the team.</i>						
Team Members	Role/ Responsibility	Email				
IZA Prince Boaz	Data preparation, PCA core implementation, GitHub repo	b.iza@alustudent.com				
Bior Aguer Kuir	Variance analysis, visualization, benchmarking, written report	b.aguerkui@alustudent.com				
Task Allocation and Tracker						
	Task Allocated	Assigned Member	Deadline	Completion Status (Not Started/In Progress/Completed/Did not participate)	Reviewed by Team? (Yes/No)	Comments
	Download EICV5 household microdata from NISR	IZA Prince Boaz	27/09/2026	Completed	Yes	Dataset meets requirements: African data, 7+ columns, missing values, non-numeric columns
	Convert data to CSV and select PCA features	IZA Prince Boaz	27/09/2026	Completed	Yes	Extracted 11 household economic and demographic columns; converted to clean CSV
	Clean missing values, encode categories, prepare data	IZA Prince Boaz	28/09/2026	Completed	Yes	Data handling completed with missing value imputation and encoding using NumPy exclusively
	Standardize data, compute covariance matrix, write notes	IZA Prince Boaz	28/09/2026	Completed	Yes	Features standardized and covariance matrix computed; covariance explanation completed
	Compute eigendecomposition and sort eigenvalues/eigenvectors	IZA Prince Boaz	29/09/2026	Completed	Yes	Eigenvalues and eigenvectors computed and sorted in descending order; verified
	Calculate explained variance and select components	Bior Aguer Kuir	29/09/2026	Completed	Yes	Explained and cumulative variance calculated; component trade-off explanation documented
	Project data onto PCs, plot results, write analysis	Bior Aguer Kuir	29/09/2026	Completed	Yes	Data projected onto principal components; before and after plots and interpretations generated

	Optimize PCA implementation and benchmark runtime	Bior Aguer Kuir	30/09/2026	Completed	Yes	Vectorized PCA optimization completed; runtime benchmarked and analyzed
	Finalize GitHub repo, export PDFs, and submit	IZA Prince Boaz	30/09/2026	Completed	Yes	Final notebook, data, and documentation committed; repo: https://github.com/boaziza/formative2-pca-pair-team-34

Meeting Attendance Log

Meeting Date	Agenda	Facilitator	Attendees	Key Discussion Points	Next Steps	
2026-09-27	Project kickoff, dataset selection, and task delegation	IZA Prince Boaz	IZA Prince Boaz, Bior Aguer Kuir	Reviewed requirements, selected EICV5 dataset, split tasks	Download microdata, export CSV, begin data prep	
2026-09-28	Data preprocessing and covariance calculation review	Bior Aguer Kuir	IZA Prince Boaz, Bior Aguer Kuir	Verified feature selection, imputation logic, and covariance formula	Finalize covariance script and start eigendecomposition	
2026-09-29	Eigendecomposition and component selection strategy	IZA Prince Boaz	IZA Prince Boaz, Bior Aguer Kuir	Confirmed eigenvalues sort order, variance threshold, and elbow plot	Code variance plot, projection functions, and explanations	
2026-09-30	PCA projection visualization and runtime optimization	Bior Aguer Kuir	IZA Prince Boaz, Bior Aguer Kuir	Reviewed PCA scatter plots and loop vs vectorized benchmarks	Run final benchmarks and document information retention/loss	
2026-09-30	Final peer review, code audit, and Canvas submission prep	IZA Prince Boaz	IZA Prince Boaz, Bior Aguer Kuir	Audited notebook for pure-NumPy compliance and reviewed repo	Push final commits and submit deliverables on Canvas	